




# Lecture 11-12: Routing



Olof Hagsand KTH CSC




# What is a router?



- Host (end-system)
  - One or many network interfaces
  - Cannot forward packets between them
- Router
  - Can forward packets between multiple interfaces
  - Forwarding on Layer 3

# What does a router do?



- Packet forwarding
- Not only IPv4:
  - IPv6, MPLS, Tunneling,...
  - (But never naming,...)
- Filter traffic
  - Access lists based on src/dst, etc.
- Metering/Shaping/Policing
  - Measuring, forming and dropping traffic
- Compute routes: build forwarding table
- In the "background": routing
- In "real-time": forwarding



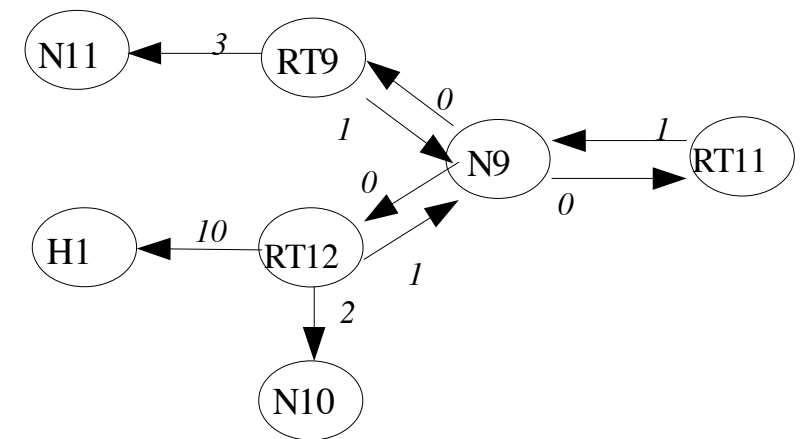

# Routing algorithms

- How does a router find a best path?
- Most solutions based on SPF (Shortest Path First) algorithms that are well known in graph theory.
  - Bellman-Ford
  - Dijkstra
- Link-State protocols (OSPF, IS-IS) use Dijkstra
- Distance-Vector protocols (RIP, IGRP, BGP) use Bellman-Ford
- Apart from that, there may also be other algorithms in
  - Multicast routing
  - Ad-hoc routing
  - Sensor networks
  - Delay-tolerant networks




# Graphs vs networks

- Algorithms are usually defined on graphs whereas protocols work on networks
- Graphs have nodes and edges whereas networks have interfaces, broadcast links, addresses, hierarchical layering, etc.
- Note the modelling of the broadcast link N9




# Shortest Path First (SPF)

- Given link metrics (weights) on each individual link
- Find the path (sequence of links) where the sum of the metrics of all links (cumulative cost) is lowest
- Equal cost multipath (ECMP): A set of paths with the least (same) cost
- What is the SPF from A to F?



# Alternative: Widest path first

- Numbers denote width: load or bandwidth
  - Available bandwidth
- It is easy to extend SPF algorithms with a widest-path computation rather than shortest path.
- What is the widest path from A -> E?



# Distance-Vector/Bellman-Ford



- Each router sends a list of distance-vectors (route with costs) to each neighbour periodically
- Every router selects the route with smallest metric (positive integer)
- The underlying algorithm is called Bellman-Ford.
- Protocols that use Bellman-Ford are called Distance-vector protocols




# Example: Distance-vector

**A:s initial state: (directly connected networks)**

Dest	Cost	NextHop
B	1	-
D	3	-

**A distributes this DV to its neighbours (B and D)**



**A receives B:s (initial) distance vector**


Dest	Cost
A	1
C	2
E	4

**A:s state after merging B:s DV:**

Dest	Cost	NextHop
B	1	-
C	3	B
D	3	-
E	5	B

**A distributes this DV to its neighbours (B and D)**

# Example: Complete and final state



	A	B	C	D	E
A		1		3	
B	1		2		4
C		2			5
D	3				6
E		4	5	6	

**Link metric matrix**

	A	B	C	D	E
A	0				
B		0			
C			0		
D				0	
E					0

**Initial state**

A's Distance-Vector

	A	B	C	D	E
A	0	1	3	3	5
B	1	0	2	4	4
C	3	2	0	6	5
D	3	4	6	0	6
E	5	4	5	6	0

**Final state**


# The algorithm

- Keep a table with an entry for each destination  $D$  in the network.
- Store the metric  $M$  (distance) and next-hop  $N$  for each  $D$  in the table.
- Periodically, send the table to all neighbors (the distance-vector).
- For each update that comes in from neighbor  $N'$  (to  $D$  with a new metric):
  - Add the cost of the link to  $N'$  to the new metric to get  $M'$ .
  - Replace the route if  $M' < M$ .
  - If  $N = N'$ , always replace the route.
- In most protocols,  $M$  is bounded, typically to 16. This upper bound is defined as unreachable(infinity).



# Going to real networks

- IP networks require destinations and nexthops (not just nodes)
  - Destinations are networks eg 192.16.32.0/24
  - Next-hops are IP addresses, eg 192.16.32.1
- Suppose the topology changes , eg routers, links crash?
  - Use timers (counters) and age the entries
  - Send updates every (e.g.) 30s
  - If you do not hear from a router in (e.g.) 180s, mark it as invalid




Dest	Cost	NextHop
1.1.1.0/24	3	-
2.2.2.0/24	1	-
3.3.3.0/24	5	2.2.2.2
4.4.4.0/24	9	1.1.1.2
5.5.5.0/24	3	2.2.2.2
6.6.6.0/24	8	2.2.2.2


**Converged routing state of A**

# D.V. Problem: Count to Infinity (Two-node instability)


Initially,  $R_1$  and  $R_2$  both have a route to  $N$  with metric 1 and 2, respectively.




The link between  $R_1$  and  $N$  fails.



Now  $R_1$  removes its route to  $N$ , by setting its metric to 16 (infinity).



Now two things can happen: Either  $R_1$  reports its route to  $R_2$ . Everything is fine.



# D.V. Problem: Count to Infinity

The other alternative is that  $R_2$ , which still has a route to N, advertises it to  $R_1$ . Now things start to go wrong: packets to N are looped until their TTL expires!




Eventually (~10-20s),  $R_1$  sends an update to  $R_2$ . The cost to N increases, but the loop remains.

Yet some time later,  $R_2$  sends an update to  $R_1$ .

...

Finally, the cost reaches infinity at 16, and N is unreachable. The loop is broken!




# Solution: Split Horizon

- *Do not send routes back over the same interface from where the route 'arrived'.*
- This helps in avoiding “mutual deception”: two routers tell each other they can reach a destination via each other.



$R_2$ , does not announce the route to N to  $R_1$  since that is where it was learnt.

Eventually,  $R_1$  reports its route to  $R_2$  and everything is fine.




# Solution: Split Horizon + Poison Reverse

- *Advertise reverse routes with a metric of 16 (i.e., unreachable).*
- Does not add information but breaks loops faster
- Adds protocol overhead



$R_2$  always announces an unreachable route to N to  $R_1$ .


Eventually,  $R_1$  reports its route to  $R_2$  and everything is fine.





# Remaining problems

- More than two routers involved in mutual deception
  - A may believe it has a route through B, B through C, and C through A
- In this case, split horizon with poison reverse does not help




# Solution: Triggered Update

- *Send out update immediately when metrics change*
- But only the changed route, not the complete table
- This may lead to a cascade of updates
  - Apply the rule above recursively!
  - Therefore, triggered updates are not allowed more often than, for example 1-5 seconds.
- A router may use triggered update only when deleting routes (16).



$R_1$  immediately announces the broken link when it happens.




# Solution: Hold Down

- *When a route is removed, no update of this route is accepted for some period of time (hold-down time)- to give everyone a chance to remove the route.*



$R_1$  ignores updates to N from  $R_2$  for some period of time.

Eventually,  $R_1$  sends the update to  $R_2$ .




# Dijkstra's shortest path first

From the link-state database, compute a shortest path delivery tree using a permanent set S and a tentative set Q:


1. Define the root of the tree: the router
2. Assign a cost of 0 to this node and make it the first permanent node.
3. Examine each neighbor node of the last permanent node.
4. Assign a cumulative cost to each node and make it tentative.
5. Among the list of tentative nodes:
  - Find the node with the smallest cumulative cost and make it permanent.
  - If a node can be reached from more than one direction, select the direction with the smallest cumulative cost.
6. Repeat steps 3 to 5 until every node is permanent.



# Example network



# Example graph



# Exercise: Dijkstra from A



Permanent set	Tentative set
A 0 -	10.0.3.0/24 1 - 10.0.1.0/24 1 -

# Exercise: Dijkstra



Permanent set

A 0 -

10.0.3.0/24 1 -

Tentative set

~~10.0.3.0/24 1 -~~

10.0.1.0/24 1 -

B 1 -



# Exercise: Dijkstra



Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -

Tentative set

10.0.1.0/24 1 -  
~~B 1 -~~  
10.0.2.0/24 2 B  
10.0.6.0/24 2 B

# Exercise: Dijkstra



## Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -

## Tentative set

~~10.0.1.0/24 1 -~~  
10.0.2.0/24 2 B  
10.0.6.0/24 2 B  
C 1 -

# Exercise: Dijkstra



## Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -  
C 1 -

## Tentative set

10.0.2.0/24 2 B  
10.0.6.0/24 2 B  
~~C 1 -~~  
10.0.2.0/24 2 C  
10.0.4.0/24 2 C

# Exercise: Dijkstra



Note: ECMP

Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -  
C 1 -  
10.0.2.0/24 2 B  
10.0.2.0/24 2 C

Tentative set

~~10.0.2.0/24 2 B~~  
10.0.6.0/24 2 B  
  
~~10.0.2.0/24 2 C~~  
10.0.4.0/24 2 C

ECMP: Equal Cost MultiPath. More than

# Exercise: Dijkstra



## Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -  
C 1 -  
10.0.2.0/24 2 B  
10.0.2.0/24 2 C  
10.0.4.0/24 2 C

## Tentative set

10.0.6.0/24 2 B

~~10.0.4.0/24 2 C~~

D 2 C

E 2 C

# Exercise: Dijkstra



## Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -  
C 1 -  
10.0.2.0/24 2 B  
10.0.2.0/24 2 C  
10.0.4.0/24 2 C  
~~E 2 C~~  
~~D 2 C~~

## Tentative set

10.0.6.0/24 2 B

~~D 2 C~~

~~E 2 C~~

10.0.5.0/24 3 C

# Exercise: Dijkstra



## Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -  
C 1 -  
10.0.2.0/24 2 B  
10.0.2.0/24 2 C  
10.0.4.0/24 2 C  
E 2 C  
D 2 C  
10.0.6.0/24 2 B

## Tentative set

~~10.0.6.0/24 2 B~~

10.0.5.0/24 3 C  
F 2 B

# Exercise: Dijkstra



## Permanent set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -  
C 1 -  
10.0.2.0/24 2 B  
10.0.2.0/24 2 C  
10.0.4.0/24 2 C  
E 2 C  
D 2 C  
10.0.6.0/24 2 B  
F 2 B

## Tentative set

10.0.5.0/24 3 C  
~~F 2 B~~  
10.0.5.0/24 3 B



# Exercise: Dijkstra (complete)



## Permanent set

## Tentative set

A 0 -  
10.0.3.0/24 1 -  
B 1 -  
10.0.1.0/24 1 -  
C 1 -  
10.0.2.0/24 2 B  
10.0.2.0/24 2 C  
10.0.4.0/24 2 C  
E 2 C  
D 2 C  
10.0.6.0/24 2 B  
F 2 B  
10.0.5.0/24 3 B  
10.0.5.0/24 3 C

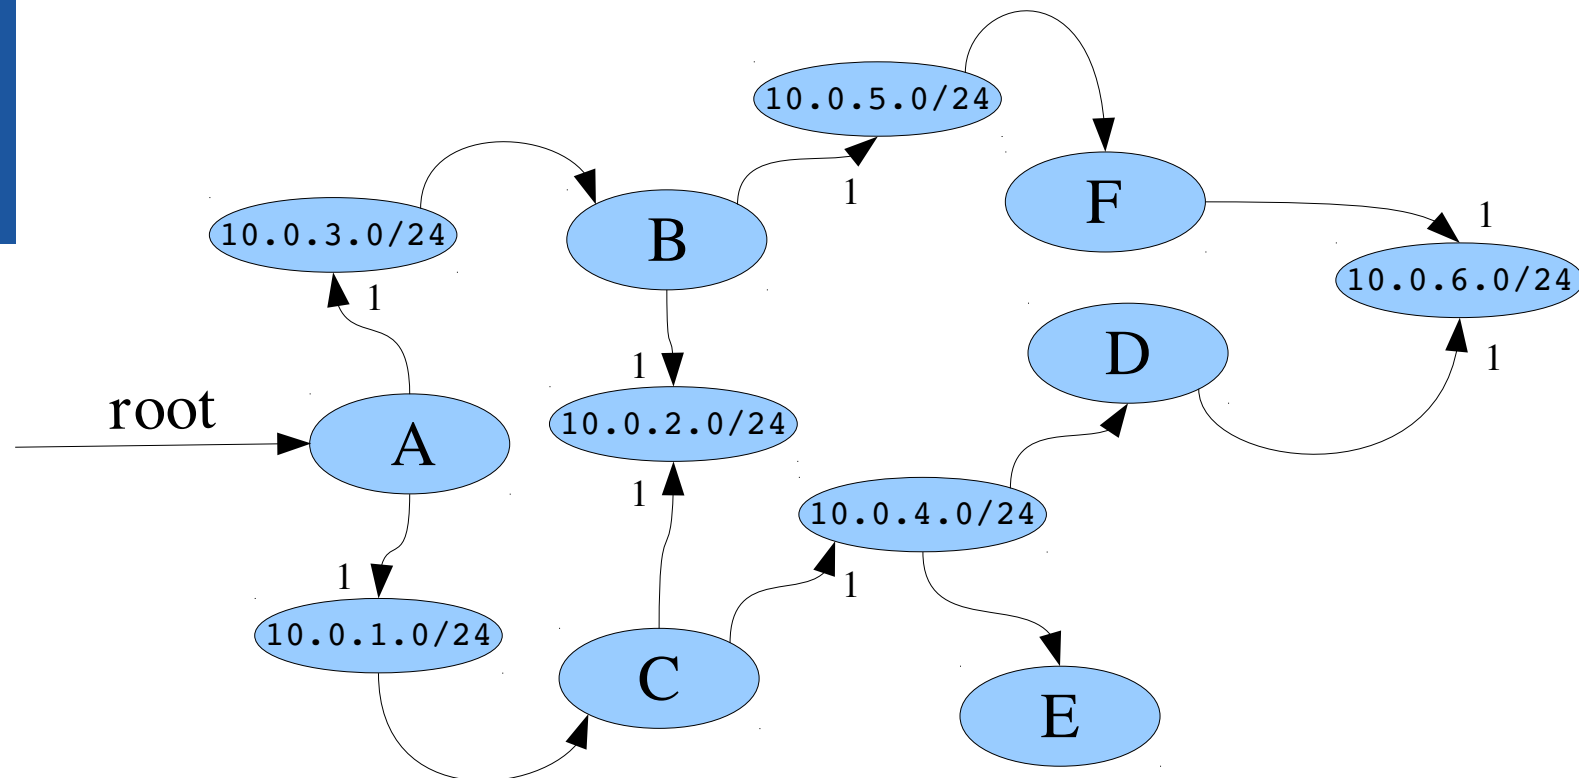
~~10.0.5.0/24 3 C~~

~~10.0.5.0/24 3 B~~

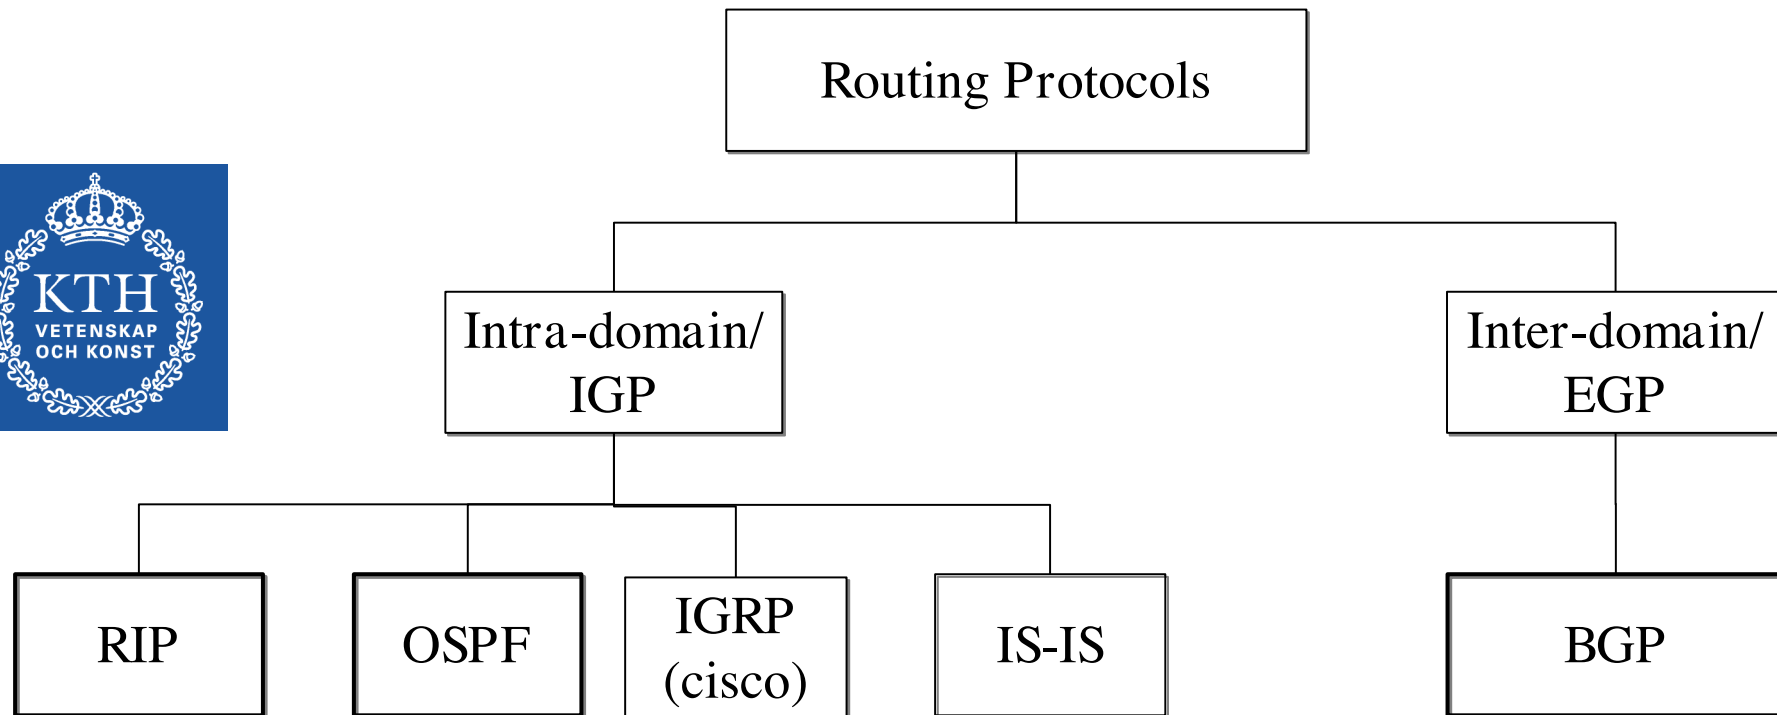
Note: ECMP

# Exercise: Dijkstra tree graph view

- Compare with table view in the previous slide
- Note the ECMP routes to 10.0.2.0/24 and 10.0.6.0/24



# Popular Unicast Routing Protocols



# Routing Information Protocol - RIP

- RIP-1 (RFC 1058), RIP-2 (RFC 2453)
- Metric is hop counts
  - 1: directly connected
  - 16: infinity
  - RIP cannot support networks with diameter  $> 14$ .
- RIP uses distance vector
- RIP messages are carried via UDP datagrams.
  - IP Multicast (RIP-2): 224.0.0.9
  - Broadcast (RIP-1)



# Disadvantages with RIP

- Slow convergence
  - Changes propagate slowly
  - Each neighbor only speaks ~every 30 seconds; information propagation time over several hops is long
- Instability
  - After a router or link failure RIP takes minutes to stabilize.
- Hops count may not be the best indication for which is the best route.
- The maximum useful metric value is 15
  - Network diameter must be less than or equal to 15.
- RIP uses lots of bandwidth
  - It sends the whole routing table in updates.



# Why would anyone use RIP?

- It is easy to implement
- It is generally available
- Implementations have been rigorously tested
- It is simple to configure.
- It has little overhead (for small networks)



# Link-state routing



- Each router spreads information about its links to its neighbours.
- This information is flooded to every router in the routing domain so that every router has knowledge of the entire network topology.
- Using Dijkstra's algorithm, the shortest path to each prefix in the network is calculated
- OSPF and IS-IS are two well-known link-state routing protocols
- OSPF is popular among organizations (KTH uses OSPF)
- IS-IS is popular among operators (SUNET uses IS-IS)

# Comparison with distance-vector

- Link-state uses a distributed database model
- Distance-vector uses a distributed processing model
- Link-state pros:
  - More functionality due to distribution of original data, no dependency on intermediate routers
    - Easier to troubleshoot
  - Fast convergence: when the network changes, new routes are computed quickly
  - Less bandwidth consuming
- Distance-vector pros:
  - Less complex - easier to implement and administrate
  - Needs less memory





# The OSPF protocol



## 1)The *hello* protocol

- Is there anybody out there?
- Detection of neighboring routers
- Election of designated routers

## 2)The *exchange* protocol

- Exchange database between neighbours

## 3)Reliable *flooding*

- When links change/age send: update to neighbours and flood *recursively*.

## 4)*Shortest path* calculation

- Dijkstra's algorithm
- Compute shortest path tree to all destinations


# The link-state

- Each router describes its environment in the form of networks (links) that it is connected to
- These link-states are the elements of the distributed database
- Fundamental task in OSPF is to distribute the link-states to all nodes in a reliable way
- Then, each node can compute Dijkstra on the *same* database
  - The result (shortest path) is consistent everywhere




# Example: OSPF link state

- Translate the network below to link states (from As point of view)




# Example: OSPF link state

- Every node creates the link-state of its connected links
- Example: A is connected to two 'transit' links, which are in-turn connected to B and C respectively
- The transit links in this case 'belongs' to A, since A is *designated router* of these sub-networks.
- A therefore distributes the three link-states in the figure:
  - One for A itself (it is connected to two transit networks)
  - Two transit links which are connected to A and B, and A and C respectively.



# Flooding link-state

- Every node distributes its link-state to all others
- Initially and after link changes (also every ~30 mins)
- Example: 'A' floods its link-state by sending it to its neighbors, who in turn distributes it to their neighbors, etc
- Flooding is made reliably and to all routers
  - No need for periodic retransmit – no waste of bandwidth
- The flooding protocol is the most complex part of OSPF (not Dijkstra!)



# Autonomous Systems (AS)

- A set of routers that has a single routing policy, that run under a single technical administration
  - A single network or group of networks
  - University, business, organization, operator
- This is viewed by the outside world as an Autonomous System
  - All interior policies, protocols, etc are hidden within the AS
- Represented in the Internet by an Autonomous System Number (ASN). 0-65535
  - Example: ASN 1653 for SUNET
- Currently, operators are switching to four-byte ASNs
  - RFC 4893: BGP Support for Four-octet AS Number Space

